

# Module 4 – Scheduling & Concurrency

The BEAM for Developers

*Erik Stenman*

# Module Overview

- Course Philosophy – The Gnome Orchestra
- Conductor / Scheduler
- The Scheduler
-  **Comprehensive Scheduling Lab**
- Summary
- Hands-On Exercises

# Course Philosophy – The Gnome Orchestra

- Extend the gnome metaphor to scheduling: imagine thousands of gnomes (processes) working in a giant workshop
- The **schedulers** are conductors who ensure every gnome gets fair time to work, no gnome monopolizes resources, and work flows smoothly between different sections of the workshop

# Conductor / Scheduler

- Each conductor (scheduler) manages their own section but coordinates with other conductors to balance the overall workload
- When a gnome finishes their task or needs to wait, the conductor smoothly switches to another gnome without missing a beat

# The Scheduler

# The Scheduler Architecture

- The BEAM scheduler is what transforms millions of lightweight processes into manageable execution on real CPU cores.
- Unlike OS threads, BEAM processes are cooperatively scheduled within the VM, giving precise control over fairness and preemption.


# 1 Scheduler / Core

BEAM runs one **scheduler thread** per CPU core by default.

Each scheduler manages its own **run queue** of processes ready to execute.

This design eliminates most locks and contention between cores.

# System Overview (4-core machine):



# Scheduler States and Transitions

Each scheduler operates in different states depending on workload:

- Active: Executing processes from the run queue (code or GC)
- Idle: No processes to run
- Load Balancing: Coordinating with other schedulers

# Active

Executing processes from the run queue

- Scheduler is busy running processes
- CPU core is fully utilized
- Processes are being preempted based on reductions

# Idle

No processes to run

- Run queue is empty
- Scheduler may sleep to save power
- Wakes up when work arrives

# Load Balancing

Coordinating with other schedulers

- Checking for work stealing opportunities
- Migrating processes between schedulers
- Balancing system-wide load


# Process State Machine

The field status in the PCB contains the process state.

A process can be in one of these states:

- free - Process has exited
- runnable - Ready to run in the ready queue
- waiting - Blocked
- running - Currently executing
- exiting - terminating
- garbing - garbage collecting
- suspended - `erlang:suspend_process/2`

# Process State Diagram



# Key Behaviors

Normal states: runnable, waiting, and running

- running → runnable: Process uses up all its reductions (4000 reductions)
- running → waiting: Process enters receive with no matching message
- waiting → runnable: Message arrives or timeout occurs
- runnable → running: Scheduler picks process from ready queue

# Garbing

Garbage collection: Process temporarily enters garbing state, saves previous state in gcstatus, then returns to previous state when GC completes.

# Suspension

Suspension: Processes can be suspended for debugging. Each `suspend_process/2` call increases the suspend count (`rcount`). Process stays suspended until `count` reaches zero via `resume_process/1` calls.

# Questions for Reflection

## Architecture Understanding

- Why does BEAM use one scheduler per CPU core instead of a global scheduler?
- What problems does per-scheduler run queues solve?
- How does this design eliminate most locking?

# Questions for Reflection

## Observation and Analysis

- Run the system analysis on your machine. How many schedulers do you have?
- What happens to run queue lengths when you create CPU-intensive processes?
- Can you predict which scheduler a new process will be assigned to?

# Questions for Reflection

## Performance Implications

- When might you want fewer schedulers than CPU cores?
- What workloads benefit most from BEAM's scheduling approach?
- How does scheduler utilization relate to system responsiveness?

# Checklist

Can you:


- explain BEAM's per-core scheduler architecture?
- explain the difference between RUNNING, WAITING, and SUSPENDED process states?
- measure scheduler utilization on your system?
- tell what factors affect run queue lengths?
- identify scheduler-related performance issues?

# Run Queues

Run queues are the heart of BEAM's scheduling system. Understanding how processes enter, wait in, and leave run queues explains most scheduling behavior you'll observe in production.

# Run Queue Structure and Priority

Each scheduler maintains multiple priority levels within its run queue:



# Priority scheduling rules

- Higher priority processes always run before lower priority
- Within same priority, processes are scheduled in FIFO order
- A running process can be preempted by higher priority processes
- Priority inversion is avoided through careful queue management

# Process Selection Algorithm

When a scheduler selects the next process to run:

- Check higher priority queues first
- Within priority level, use FIFO
- Consider process "locality" (recently run processes may have better cache)
- Apply fairness constraints (prevent starvation)

## Speaker notes

The low priority does not have its own queue but is picked every X turn in the timing wheel.

# Message-Driven Scheduling

Process transitions between states are primarily driven  
by message passing:

# Process becomes runnable when:

- A message arrives in its mailbox
- A timeout expires
- An I/O operation completes
- It's newly spawned

# Process becomes suspended when:

- It executes `receive` with no matching messages
- It calls a blocking operation (Send to busy port)
- It explicitly yields

# Timing Wheels

## How They Work

Timers are handled in the VM by a timing wheel

An array of time slots that wraps around.

Prior to Erlang 18, the timing wheel was a global resource with potential contention for the write lock when many processes inserted timers.

## Speaker notes

The timing wheel is a circular buffer mechanism that efficiently handles timeout operations for receive statements and other timer-based operations. For timeouts longer than the wheel size, the count field tracks how many complete rotations are needed before the timer fires.

# Configuration


- Default size (TIW\_SIZE): 65536 slots (or 8192 slots for small memory footprint builds)
- Current time pointer: tiw\_pos - index into the array indicating current time

# Timer Insertion Algorithm

When a timer is inserted with timeout T:

- `tiw_pos` points to current time slot
- Timer is placed at slot:  $(\text{tiw\_pos} + T) \% \text{TIW\_SIZE}$
- Array wraps around from slot 65535 back to slot 0
- Multiple timers in the same slot are linked together in a linked list


# Timing Wheel Diagram



# ErlTimer Structure

Each timer in the timing wheel is represented by an ErlTimer struct:

```
typedef struct erl_timer {  
    struct erl_timer* next;      /* next entry tiw slot */  
    struct erl_timer* prev;     /* prev entry tiw slot */  
    Uint slot;                 /* slot in timer wheel */  
    Uint count;                /* loops remaining */  
    int active;                /* 1=activ, 0=inactiv */  
    void (*timeout)(void*);  
    void (*cancel)(void*);  
    void* arg;  
} ErlTimer;
```



## Speaker notes

- next/prev: Linked list pointers for chaining multiple timers in same slot
- slot: Which timing wheel slot contains this timer
- count: For timeouts > TIW\_SIZE, this tracks remaining wheel rotations ( $T/TIW\_SIZE$ )
- active: Whether timer is currently active
- timeout: Function called when timer expires
- cancel: Optional function called when timer is cancelled
- arg: Argument passed to timeout/cancel functions

# Run Queue Dynamics and Bottlenecks

Healthy run queue characteristics:

- Most queues empty or with 1-2 processes
- Processes don't stay in queue long
- Even distribution across schedulers

# Run Queue Dynamics and Bottlenecks

Problem indicators:

- Consistently long run queues (>10-20 processes)
- Uneven distribution across schedulers
- High priority processes stuck behind normal priority

# Questions

Priority and Fairness:

- What happens if you have too many high-priority processes?
- How does BEAM prevent priority inversion?
- When should you use non-normal priorities?

# Questions

Run Queue Health:

- What run queue lengths indicate problems?
- How can you detect scheduler imbalance?
- What causes processes to accumulate in run queues?

# Questions

Message-Driven Behavior:

- Why do suspended processes not consume scheduler time?
- How does `receive` with timeouts affect scheduling?
- What makes a process transition from SUSPENDED to WAITING?

# Checklist

- Do you understand the run queue priority structure?
- Can you explain when processes transition between states?
- Can you monitor and assess run queue health?
- Do you know when to use different process priorities?
- Can you predict how message patterns affect scheduling?

# Reductions and Preemption

The reduction system is BEAM's mechanism for fair preemption. It ensures no process can monopolize a scheduler thread, making the system responsive even under heavy computational load.

# What Are Reductions?

A reduction is BEAM's unit of work measurement.  
Every operation that can potentially consume  
significant time costs reductions:

- Function calls
- Message sends
- Pattern matching
- Arithmetic operations
- Memory allocation
- Built-in function calls (BIFs)

# The Preemption Mechanism

Each process gets a **reduction budget** when scheduled (typically 2000 reductions). When the budget is exhausted:

- Process is preempted (removed from CPU)
- Process goes back to run queue (if still runnable)
- Scheduler picks next process
- Process gets new budget when rescheduled

# Reduction Counting in Practice

```
measure(F) ->
    {reductions, B} = process_info(self(), reductions),
    {_, A} = process_info(self(), reductions),
    Owerhead = A - B,
    {_, Before} = process_info(self(), reductions),
    Result = F(),
    {_, After} = process_info(self(), reductions),

    io:format("Reductions used: ~p~n",
              [(After - Before) - Owerhead]),
    Result.
```

# Reduction Counting in Practice


```
%% Simple arithmetic
measure(fun() -> 1 + 1 + 1 + 1 + 1 end).

%% Function calls
measure(fun() -> lists:sum([1, 2, 3, 4, 5]) end).

%% List processing
measure(fun() ->
    lists:map(fun(X) -> X end, lists:seq(1, 100))
end).

%% Message sending
measure(fun() -> self() ! test end).

%% Pattern matching
measure(fun() ->
    case {1, 2, 3} of {A, B, C} -> A end
end).
```



# Preemption and Latency

Understanding preemption helps predict system latency:

```
defmodule LatencyAnalysis do
  def analyze_preemption_impact do
    # Simulate competing processes with different characteristics
    # CPU-intensive process (uses full reduction budget)
    cpu_intensive = spawn(fn ->
      cpu_heavy_work()
    end)

    # Interactive process (frequent yields)
```

```
interactive = spawn(fn ->  
    interactive_work()  
end)
```

# Questions on Reductions

- Reduction Understanding:
- Why does BEAM count reductions instead of using time-based preemption?
- Which operations are "expensive" in terms of reductions?
- How do reductions relate to actual CPU time?

# Questions on Reductions

- Preemption Behavior:
- When exactly does preemption occur?
- Can a process be preempted in the middle of a function call?
- How does preemption affect message processing latency?

# Questions on Reductions

- **Performance Implications:**
- When should you manually yield control?
- How do reduction costs affect system design?
- What happens if processes never yield?

# Checklist

- Do you understand what reductions measure?
- Can you predict which operations will be expensive?
- Do you know when processes get preempted?
- Can you measure reduction usage in your code?
- Do you understand when to yield manually?

# SMP and Load Balancing

BEAM's Symmetric Multi-Processing (SMP) support enables true parallelism across CPU cores.

# Key principles

- Each scheduler operates independently
- Minimal locking between schedulers
- Work stealing for load balancing
- Process migration for optimization

# Work Stealing and Migration

When schedulers become imbalanced, BEAM uses  
**work stealing** to redistribute load:

# Work stealing triggers:

- One scheduler has empty run queue while others are loaded
- Significant imbalance detected during load checks
- Scheduler utilization falls below threshold

# Migration strategies:

- Steal from heavily loaded schedulers
- Prefer recently spawned processes (less cache pollution)
- Respect process affinity when possible
- Avoid migrating message-heavy processes

# Process Affinity and Locality

BEAM tries to maintain **process affinity** to improve cache performance:

# Affinity benefits:

- Better CPU cache utilization
- Reduced memory latency
- Less cross-core communication

# When affinity is broken:

- Work stealing for load balancing
- Process migration for resource access
- Explicit scheduler assignment

# VM flags for scheduler tuning:

```
# Set number of schedulers
erl +S 4:4          # 4 schedulers, 4 online

# Scheduler binding
erl +sbt db          # Default binding (recommended)
erl +sbt u           # Unbound
erl +sbt ns          # No spread

erl +swt very_low|low|medium|high|very_high -
```

Speaker notes

Does not work on OsX

# Questions on Run Queues

- SMP Architecture:
- Why does BEAM use separate run queues per scheduler?
- What are the trade-offs of work stealing?
- How does process migration affect performance?

# Questions on Run Queues

- Load Balancing:
- When does work stealing occur?
- What factors determine if a process should be migrated?
- How can you detect scheduler imbalance?

# Questions on Run Queues

- Configuration and Tuning:
- When might you want fewer schedulers than CPU cores?
- How does scheduler binding affect performance?
- What workloads benefit from SMP scheduler tuning?

# Checklist

- Do you understand BEAM's SMP architecture?
- Can you explain work stealing and migration?
- Do you know how to measure scheduler utilization?
- Can you identify load balancing issues?
- Do you understand scheduler configuration options?



# Comprehensive Scheduling Lab

# Lab 1: Complete Scheduler Analysis

Build a comprehensive scheduler monitoring tool:

```
defmodule SchedulerAnalyzer do
  def full_system_analysis(duration_seconds \\ 30) do
    IO.puts("==> Complete Scheduler Analysis (#"
    {duration_seconds}s) ==>")

    # Initialize monitoring
    :erlang.system_flag(:scheduler_wall_time, true)
    :erlang.statistics(:scheduler_wall_time) # Reset
    baseline
```

```
# Start workload
workload_pids = create_comprehensive_workload()

# Collect data over time
```

# Lab 2: Scheduler Tuning Experiment

Test different scheduler configurations:

```
defmodule SchedulerTuningLab do
  def test_different_configurations do
    IO.puts("==> Scheduler Configuration Testing ==>")
    # Test 1: Priority impact
    test_priority_effects()
    :timer.sleep(2000)
    # Test 2: Process affinity
```

```
test_process_affinity()
```

```
:timer.sleep(2000)
```

- **System Design:**
- What scheduler metrics would you monitor in production?
- How do you balance throughput vs latency in process design?

- **Performance Analysis:**
- A process is using 10x more reductions than expected. How do you investigate?
- Your schedulers show 90% utilization but response times are poor. Why?
- How do you detect if work stealing is happening too frequently?

- **Tuning Scenarios:**
- When would you use fewer schedulers than CPU cores?
- How do you optimize for batch processing vs interactive workloads?
- What scheduler flags would you tune for a low-latency trading system?

# Summary

# Key Takeaways

- **Scheduler Architecture:** BEAM's per-core design eliminates most locking
- **Fair Scheduling:** Reduction-based preemption ensures responsiveness
- **Load Balancing:** Work stealing maintains balance across cores
- **Configurability:** Multiple tuning options for different workload patterns

# Production Considerations

**Monitoring:** Track scheduler utilization, run queue lengths, reduction rates

**Alerting:** Set up alerts for persistent queue imbalances or high utilization

**Capacity Planning:** Understand how your workload scales with scheduler count

# Interactive Exercises

Open the LiveBook:

module-4-exercises.livemd

20 minutes of hands-on practice with:

- Scheduler architecture and configuration analysis
- Process state transitions and priority effects
- Reduction measurement and preemption behavior
- SMP load balancing and work stealing

## Speaker notes

The exercises are in an interactive LiveBook notebook. Students should open `module-4-exercises.livemd` in LiveBook to work through the hands-on exercises.

Each exercise has runnable code cells and discussion questions.

# Final Checklist

Can you explain:

- BEAM's scheduling architecture end-to-end?
- how reductions drive preemption?
- how to analyze scheduler utilization and load balance?
- how to tune scheduler behavior?
- the trade-offs between different configurations?

# Hands-On Exercises

# Question 1

How many scheduler threads does BEAM create by default?

- One per logical CPU (including hyperthreading)
- One per physical CPU core only
- Always 4 schedulers regardless of hardware
- As many as needed based on process count

# Question 2

What happens when a process exhausts its reduction budget?

- It is preempted and moved to the back of its run queue
- It crashes with a reduction\_limit error
- It continues running until it calls receive
- The scheduler picks the next process to run

# Question 3

Which process states do NOT consume scheduler time?

- SUSPENDED (blocked in receive)
- RUNNING
- WAITING (in run queue but not executing)
- GARBING (performing GC)

# Question 4

What is work stealing in BEAM's scheduler?

- Processes stealing CPU time from other processes
- Idle schedulers taking processes from busy schedulers' run queues
- High priority processes preempting low priority ones
- Load balancing mechanism to distribute work across cores

# Question 5

Why does BEAM use reduction-based scheduling instead of time-based?

- Ensures processes yield at safe, known execution points
- Avoids complexity of interrupting arbitrary VM state
- Time-based scheduling is impossible in a VM
- Reductions are faster to count than time

# Question 6

What is process priority's effect on scheduling?

- Higher priority processes run before lower priority ones
- Priority determines reduction budget size
- Within same priority, FIFO order is used
- Priority only affects process spawn order

# Question 7

When should you manually yield with  
`timer:sleep(0)`?

- When processing large batches to allow other processes to run
- Never, the scheduler handles everything automatically
- In long-running loops that don't hit receive
- After every function call for fairness

# Question 8

What does a consistently high run queue length indicate?

- Efficient use of CPU resources
- More runnable processes than scheduler capacity
- Possible performance bottleneck
- Normal operation under any load

# Question 9

How does process migration affect performance?

- Reduces CPU cache locality when process moves to new scheduler
- Always improves performance through better load balance
- Trades cache performance for load balancing benefits
- Has no performance impact

# Question 10

What operations automatically yield control before exhausting reductions?

- `receive` with no matching messages
- Long-running BIFs like file I/O
- Arithmetic operations
- Function calls

# Navigation

[← Previous Module: Memory & Garbage Collection](#) |  
[Back to Course Structure](#) | [Next Module: Debugging & Observability →](#)